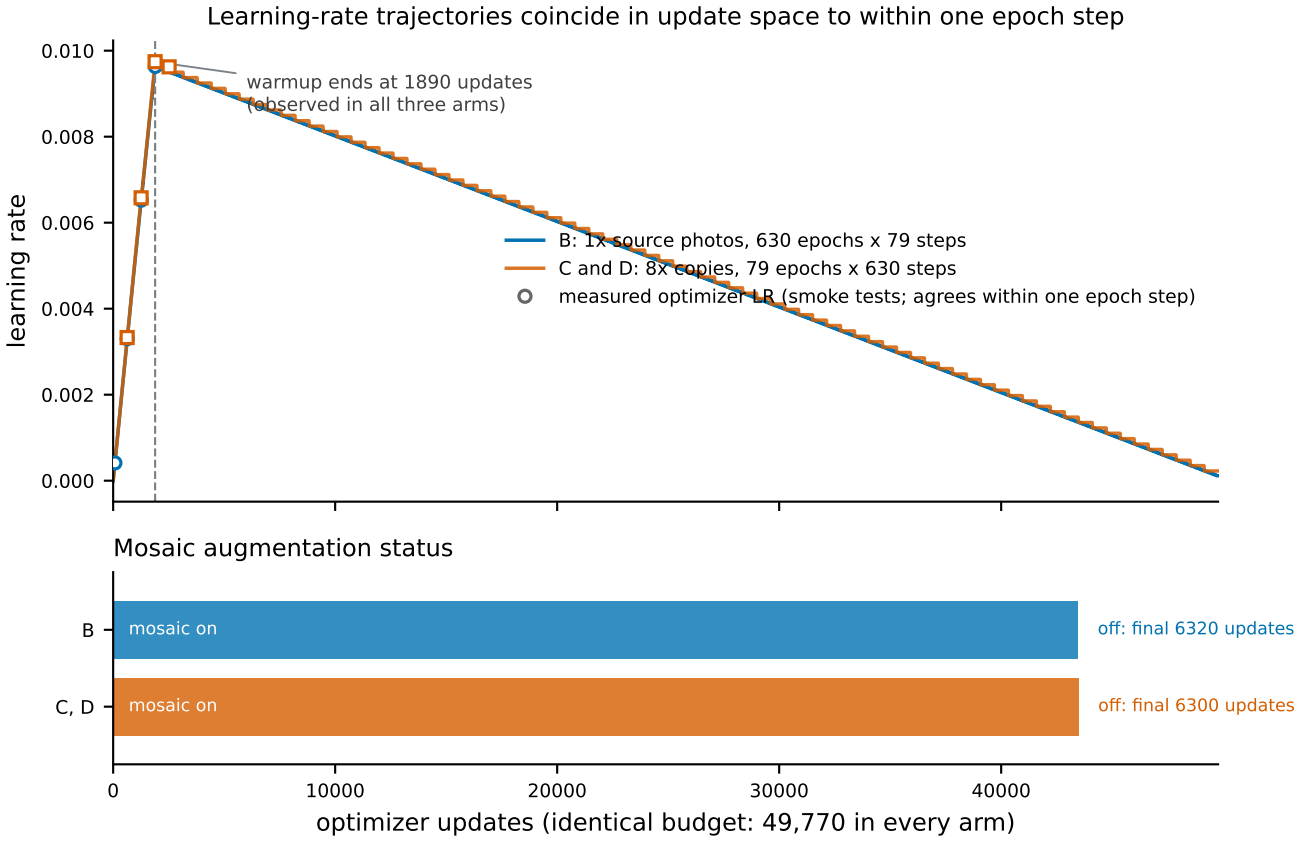


Curve: analytic reproduction of the trainer's epoch-quantised LambdaLR, verified against logged optimizer state.
Agreement is within 0.05% of $lr_0$ through warmup; afterwards within one epoch step (0.16% of $lr_0$ for B, 1.25% for C and D),
which is the sampling resolution, not a model error. Inter-arm divergence is therefore resolved only to one epoch step.